Speaking GitHub Copilot's Language

Context engineering for gigantic codebases

Tim Warner

Principal Author | IT Ops | Pluralsight

TechTrainerTim.com

Don't bring the monorepo to the model.

Bring the model to the right neighborhood.

Context is a finite budget, not a landfill.

Exercise files

github.com/timothywarner-org/eShop

Instead of stuffing context into your LLM...

...sculpt your LLM around your context

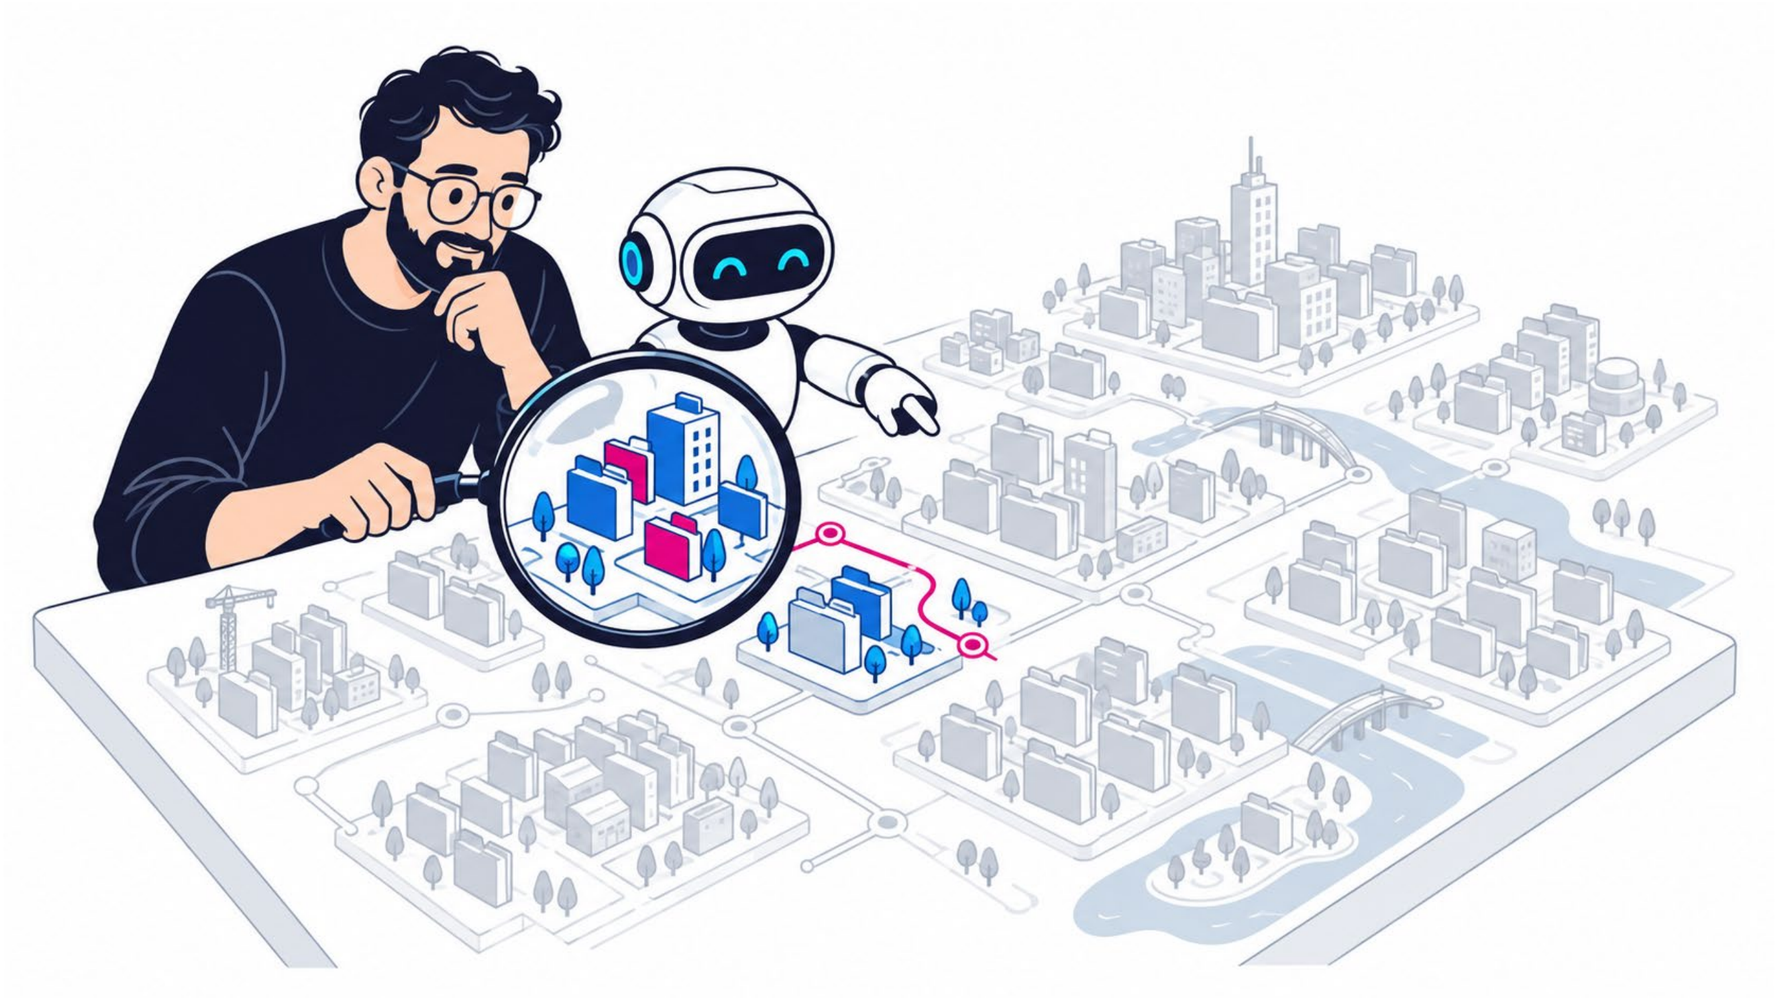

The session holds the state

Stateful Copilot session

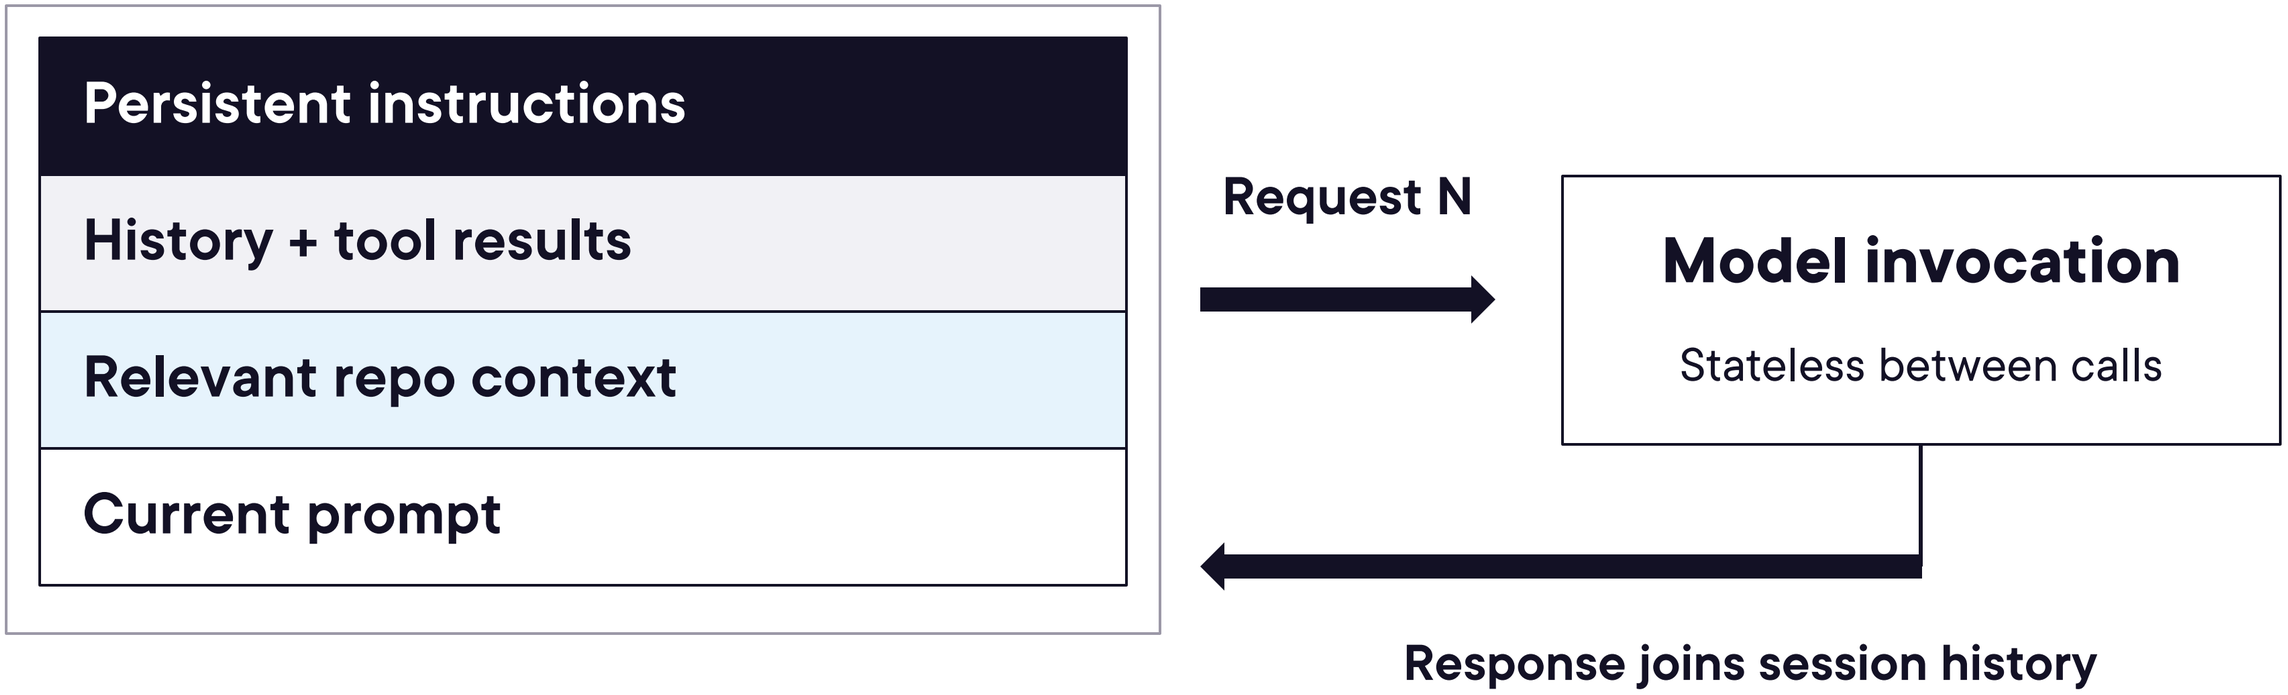

Next call: assemble a fresh request from relevant state.

System instructions, tool definitions, and editor context also use the budget.

The payload grows. The budget doesn't

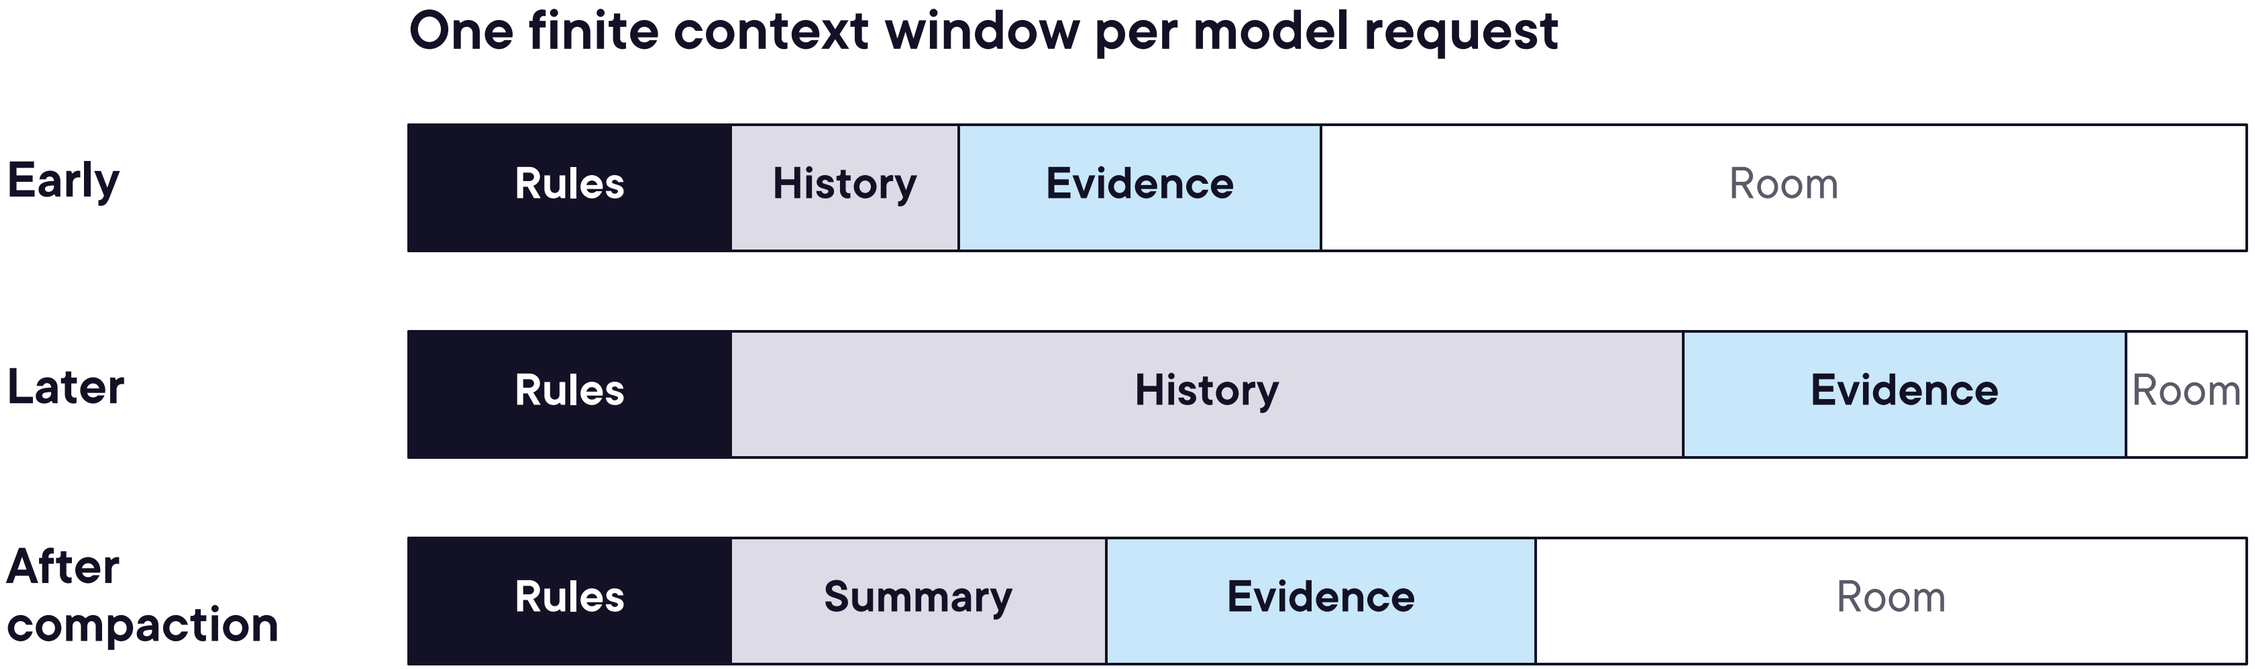

Keep the goal and decisions. Re-read exact evidence when needed.

Same task: compact history. Different task: start a fresh session.

Conceptual proportions. Compaction summarizes older history and can omit detail.

Standing guidance applies across the repo

.github/copilot-instructions.md

Short, stable, actionable. Always-on for Copilot Chat.

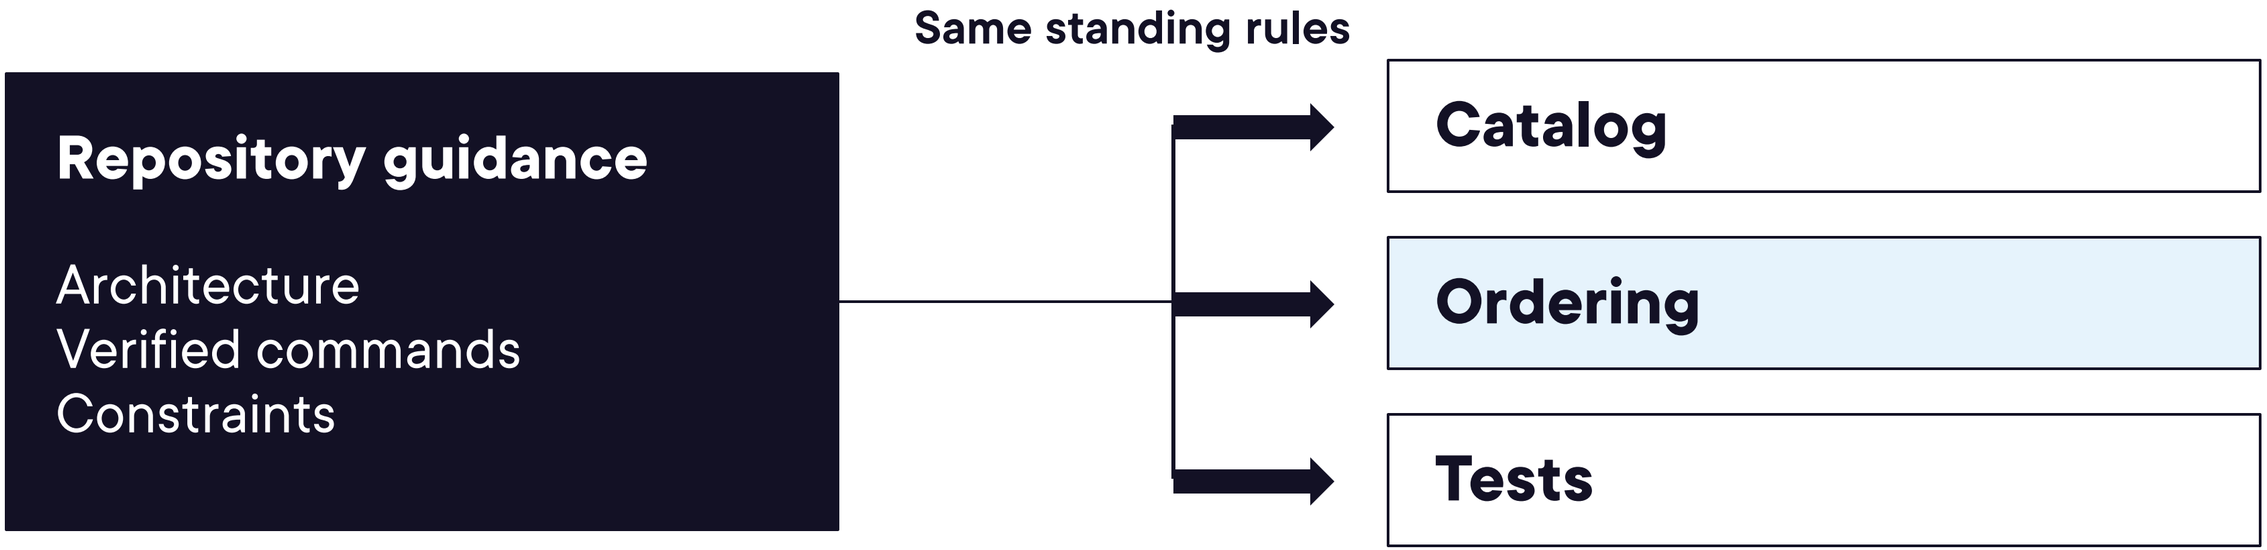

Copilot CLI: /init

Creates or updates the file. Review it.

One source of truth per rule

Mechanism	Scope
copilot-instructions.md	Repository
*.instructions.md	Matching paths
AGENTS.md	Optional portability

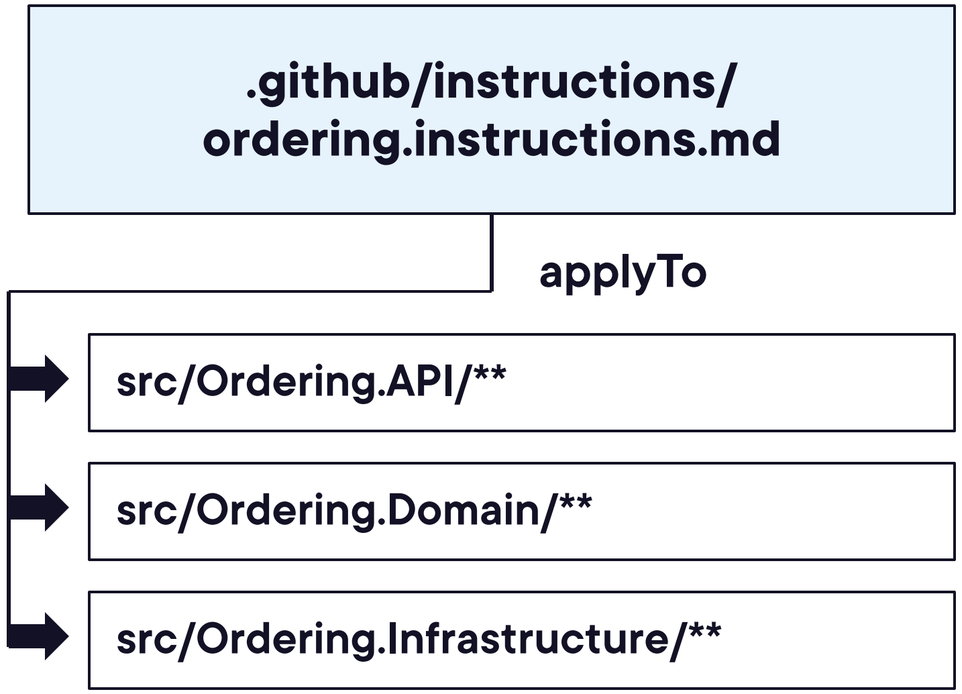

Directory scoping with AGENTS.md depends on the client.

In VS Code, nested AGENTS.md support is experimental.

Add a rule where it belongs. Avoid mirrored instructions.

Retrieval selects the working set

Task: trace the buyer and payment IDs

Semantic index: find code by meaning. Source: check the claim.

Check index status. Search and file reads can still work without it.

Now: one eShop flow in VS Code

C:\github\eshop-1

01	02	03	04	05
Orient	**Trace**	**Focus**	**Change**	**Verify**
Map + rules	Order placement	Buyer/payment IDs	One scoped edit	Source + diff

Where do Buyer.Id and Payment.Id reach the order?

Cite the path. Check the claim. Preserve the useful rule.

Thank you for your time!

Tim Warner | TechTrainerTim.com